AI the vL Way
Less hype. More help.

Does the folder earn its place?

A folder of instructions is meant to help an AI answer questions about a set of documents. This test checked whether it does.

No. The folder got the AI to its answers in half the steps, but its answers were marked right less often than with the documents alone.

The experiment

The same 10 questions, asked of the same AI, in three set-ups. Each set-up was asked every question 3 times.

Nothing

No documents, no folder. The AI only has the question.

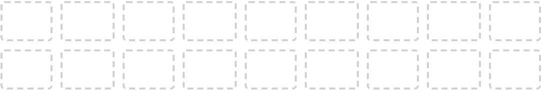

No instructions

0
marked right, of 30

2.00
steps per question

Documents alone

18 documents to search. No instructions.

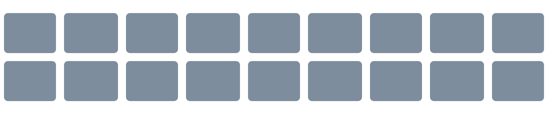

No instructions

30
marked right, of 30

5.03
steps per question

Documents with the folder

The same 18 documents, plus 4 instruction files read first.

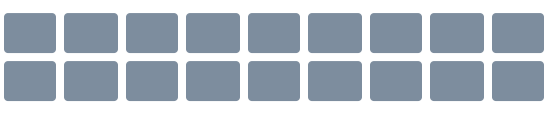

map · role · rules · examples

23 to 28
marked right, of 30

2.37
steps per question

The design

What makes it an experiment and not a demonstration.

The one thing changed	Whether the folder's 4 instruction files are loaded before the question.	**Kept the same**	The AI model and its settings, the 18 documents, the 10 questions, the tools (search, list, open), and a fresh session for every question.
What was measured	Steps taken, answers marked right, text put in front of the AI, tokens and seconds.	**The controls**	Nothing at all. The documents alone. The folder present but not loaded. A folder with its map made wrong on purpose.
The sample	10 questions, asked 3 times each, in 19 set-ups: 570 sessions. Then 10 more questions held back as a check.	**The decision rule**	A difference counts only if it shows on 9 or 10 of the 10 questions. Written down before any document existed.
Keeping bias out	Another AI wrote the questions and never saw the folder. A third AI checked the answers. The exam was drawn at random. The order of sessions was random. Graders were not told which set-up wrote an answer.	**Keeping it honest**	The sentences that could be published were written first. Every session is kept. The results are built from the raw logs by a script that anyone can run again.

How it was run

Each step was committed in this order. The rule, the questions and the folder were all fixed before the run.

1. **Rule fixed** How the folder would be judged, written before anything else.
2. **Documents** ChatGPT wrote 18 documents for two made-up organisations.
3. **Questions** ChatGPT wrote 20, with answers. It never saw the folder.
4. **Answers checked** Gemini answered all 20 with no answer key. No figure differed.
5. **Random draw** A script picked the 10 exam questions. Nobody chose.
6. **570 sessions** Claude answered in 19 set-ups, from one command.
7. **Graded 3 times** Twice by Claude, once by Gemini, set-up hidden.
8. **Results built** By script, from the raw logs. All of it is in the repo.

The results

🟧 Documents with the folder ⬛ Documents alone

Steps to reach an answer

Searches and files opened, per question. Fewer with the folder on 10 of 10.

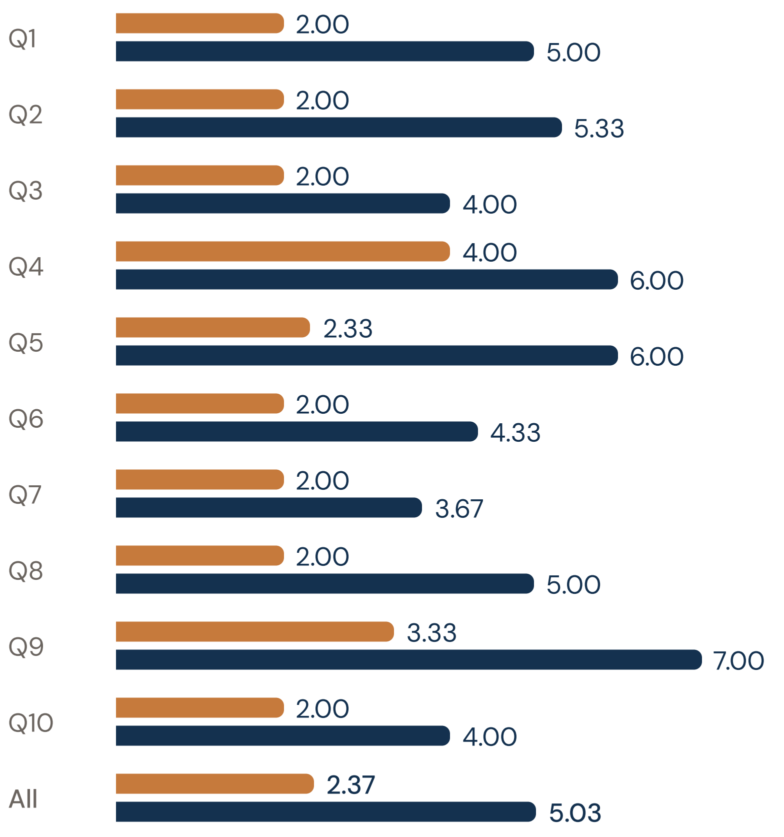

Answers marked right, of 30

Each of the three gradings.

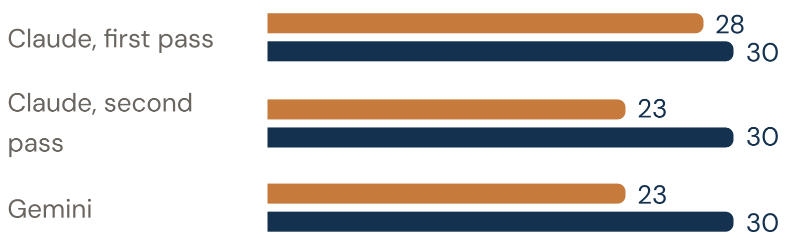

Text put in front of the AI

Characters per question. The lighter part is the folder's own files, read every time.

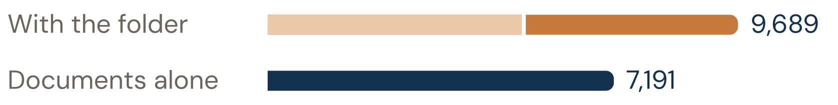

Cost per question

Tokens are the units an AI is billed in.

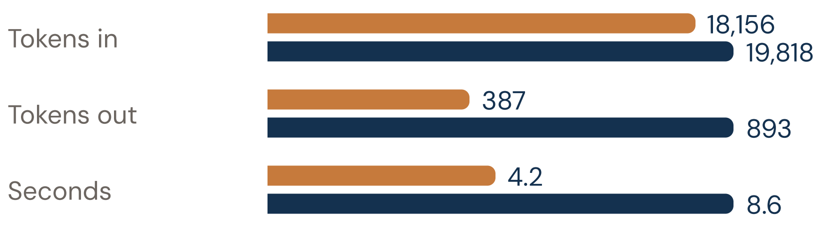

The rule, fixed before the run

The folder earns its place only if all three hold.

✓ **Steps clearly fall**
Met. Fewer on 10 of 10 questions.

✗ **Right answers do not fall**
Not met. 23 to 28 against 30.

✗ **Text does not rise**
Not met. 9,689 against 7,191.

Why the folder lost marks

The folder tells the AI to answer what was asked and then stop. Its answers were right on the main point and less than half the length. The answer key also wanted the detail a longer answer includes, such as how far over a limit something was.

What held up

- The map does the work on steps: 2.50 with the map alone, 6.07 with nothing loaded.
- A re-run from a fresh copy gave the same result: 2.50 against 5.10.
- On 10 questions kept aside: fewer steps on all 10, and 30 of 30 marked right in both set-ups. The loss of marks did not repeat.

What this does not show

- One AI model, made-up documents, ten questions, one day.
- The graders could tell the two set-ups apart from the writing.
- Nobody else has run it on another machine yet.

Where to look. `START-HERE.md` is the guide. `results.md` has every figure, tie and loss.

Weekly Comp #14, The Control. Run on 12 October 2026. AI the vL Way. Less hype. More help.