The Control
Weekly Comp #14

Does the folder earn its place?

A folder of instructions is meant to help an AI answer questions about a set of documents. This test checked whether it does.

No. The folder got the AI to its answers in half the steps, but its answers were marked right less often than with the documents alone.

The experiment

The same 10 questions, asked of the same AI, in three set-ups. Each set-up was asked every question 3 times.

Nothing
No documents, no folder. The AI only has the question.
No instructions
0 marked right, of 30
2.00 steps per question

Documents alone
18 documents to search. No instructions.
No instructions
30 marked right, of 30
5.03 steps per question

Documents with the folder
The same 18 documents, plus 4 instruction files read first.
map role rules examples
23 to 28 marked right, of 30
2.37 steps per question

The brief asks for the first comparison: the folder against nothing. The folder wins that one, because the documents do the work. The middle set-up is the control that matters. It shows what the 4 instruction files add on top of the documents.

The results

🟧 Documents with the folder ⬛ Documents alone

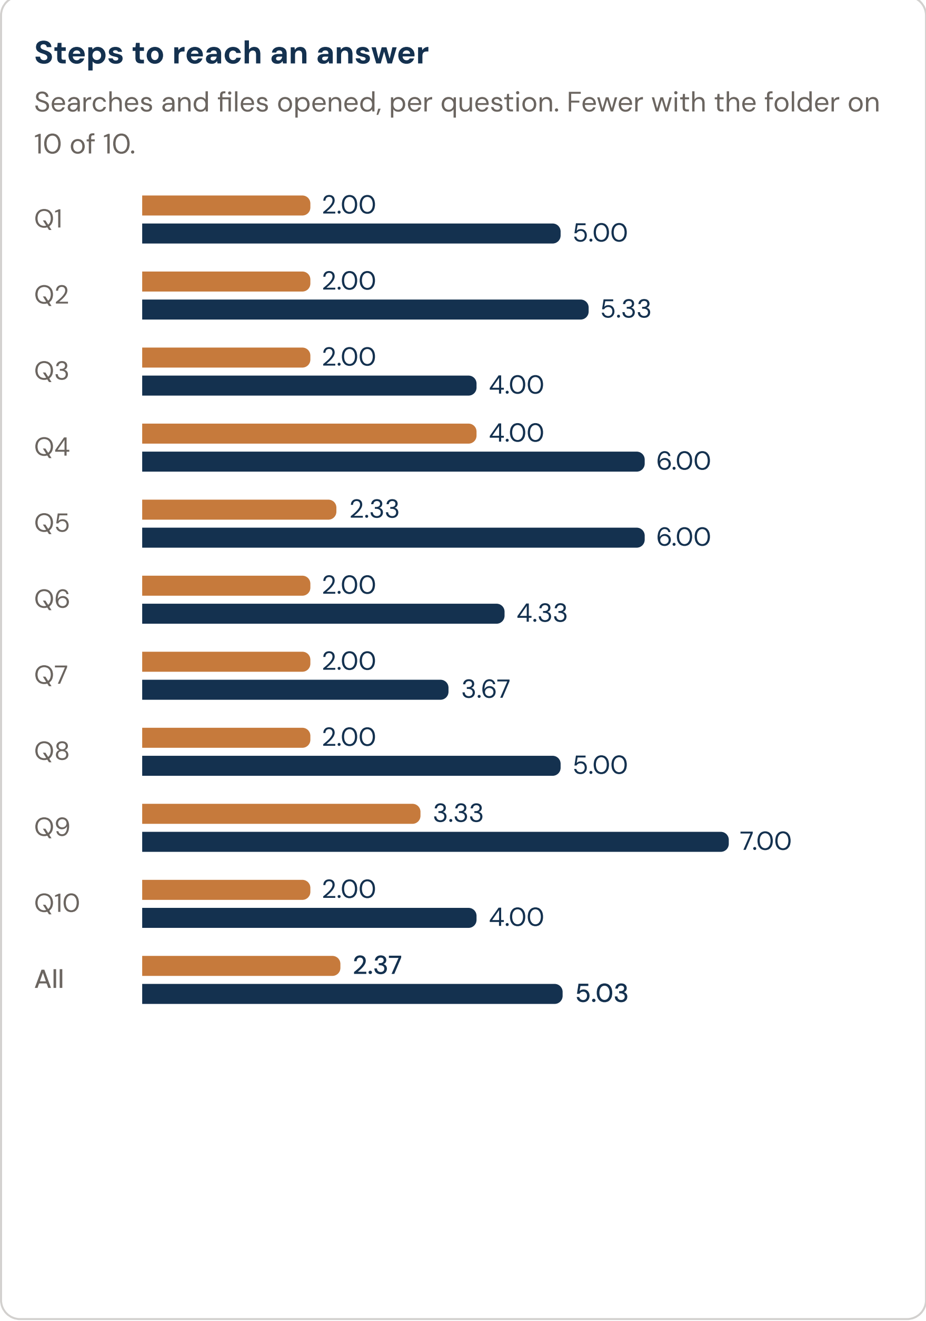
Steps to reach an answer

Searches and files opened, per question. Fewer with the folder on 10 of 10.

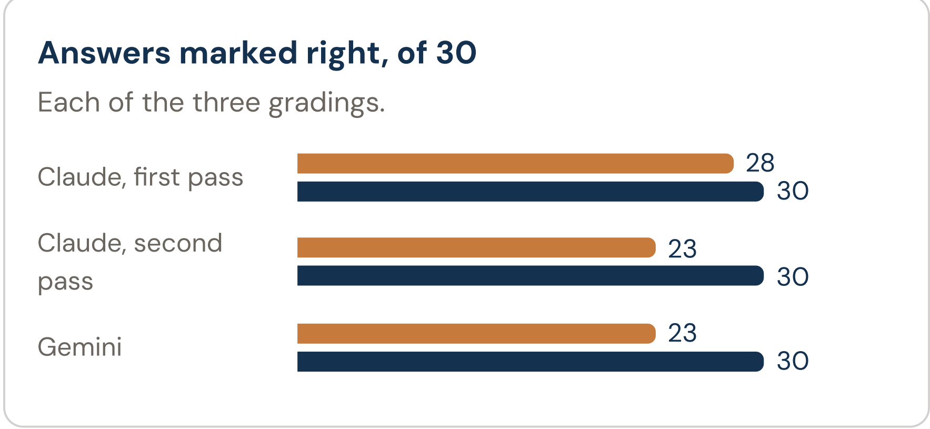
Answers marked right, of 30

Each of the three gradings.

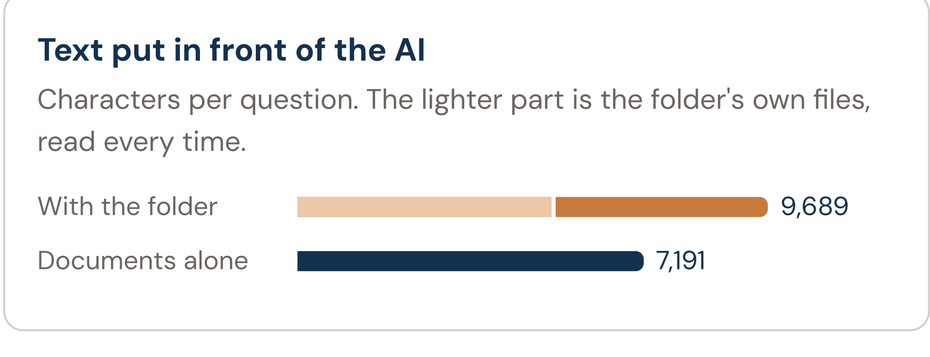
Text put in front of the AI

Characters per question. The lighter part is the folder's own files, read every time.

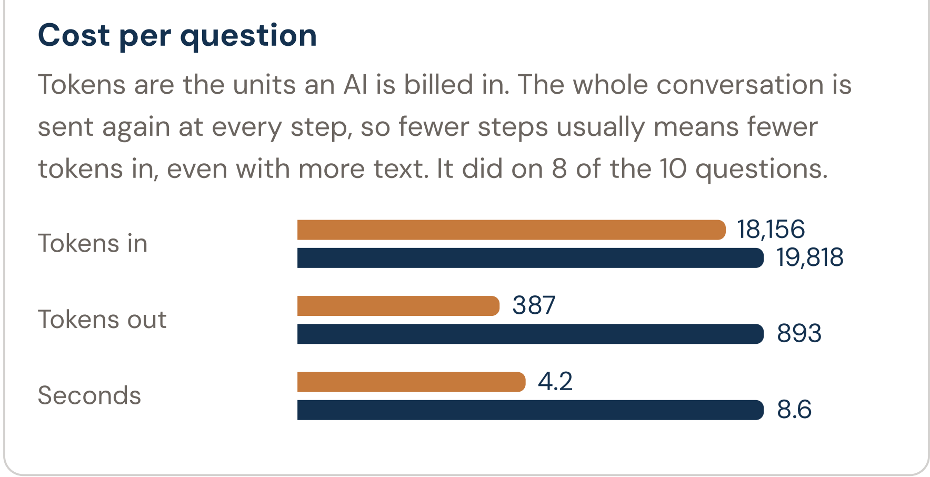
Cost per question

Tokens are the units an AI is billed in. The whole conversation is sent again at every step, so fewer steps usually means fewer tokens in, even with more text. It did on 8 of the 10 questions.

The rule, fixed before the run

The folder earns its place only if all three hold.

✓ **Steps clearly fall**
Met. Fewer on 10 of 10 questions.

⊗ **Right answers do not fall**
Not met. 23 to 28 against 30.

⊗ **Text does not rise**
Not met. 9,689 against 7,191.

Why the folder lost marks

Our reading of the sealed run, which is not proved. The folder's rules and examples push for short answers. Its answers were right on the main point and less than half the length. The answer key also wanted the detail a longer answer includes, such as how far over a limit something was.

Which part of the folder did what

The folder does two jobs: the map helps the AI find the answer, and the rules and examples shape how it writes the answer. The four instruction files were loaded in every mix, 16 set-ups with the full folder as one of them, so the two jobs can be told apart. Here is each file loaded alone.

Loaded before the question	Steps per question	Answers with every fact the key asks for, of 30
Nothing loaded	6.07	29
The map alone map	2.50	28
The role alone role	4.60	29
The rules alone rules	3.70	22
The examples alone examples	4.60	24
All four map role rules examples	2.37	21

The map does the finding: alone it gets nearly all of the saving in steps. The facts went with the other job. The 4 set-ups with neither the rules nor the examples loaded gave every fact in 28 or 29 of 30 answers. All 12 with either one loaded gave every fact in 21 to 25. The documents alone: 30.

The last column is counted by a script that searches each answer for the facts the answer key asks for, so it gives the same number every time. It was written after the run and is not a grade. In these rows the four instruction files sat in the folder even when they were not loaded, so "nothing loaded" is 6.07 steps and not the 5.03 of the documents alone.

The follow-up: one line, tested on a copy

After the verdict, one line was added to a copy of the folder's rules: "Say what would otherwise apply, and the difference." It was run alone, and then the next day beside the unchanged folder, both graded on one sheet.

✓ **Run alone, it scored 30 of 30**
On all three gradings. We first read that as the lost marks coming back.

⊗ **Side by side, it did no better**
Unchanged folder 30, 30 and 30 of 30. Changed folder 29, 29 and 30. Steps the same: 2.40 each.

⊗ **The text is still more**
9,643 against 7,395 side by side. By the same rule, still not earned.

So the line is not shown to help: it made answers about 19 words longer and did not raise the marks. In that last run the unchanged folder lost none: it left out the wanted detail less often than in the first run, and the graders passed the answers that still did. Across three batches the same folder has been graded 23 to 28, then 24 to 30, then 30 of 30. Counted by the script it gave every fact in 21, 24 and 26 of 30, with the documents alone at 30 each time. The first verdict stands.

What held up

- Most of the saving on steps comes from the map: 2.50 with the map alone, 6.07 with nothing loaded.
- A person with no part in the build re-ran it on their own machine and got the same result on steps: 2.40 against 5.60. On a question they wrote themselves, the folder gave the right figure and the documents alone gave an out-of-date one (asked once).
- Three re-runs, graded afterwards with the set-up hidden: the documents alone were right every time. By 2 of 3 gradings the folder lost marks, only on the same 3 questions.
- On 10 questions kept aside: fewer steps on all 10, and 30 of 30 marked right in both set-ups. The loss of marks did not repeat.

What this does not show

- One AI model, made-up documents, ten questions, one main run.
- The graders could tell the two set-ups apart from the writing.
- How many marks the folder loses depends on who grades: 1, 0 and 6 of 30 across three gradings of the re-runs.

The design

How the test was built. What makes it an experiment and not a demonstration.

The one thing changed	Whether the folder's 4 instruction files are loaded before the question.	**Kept the same**	The AI model and its settings, the 18 documents, the 10 questions, the tools (search, list, open), and a fresh session for every question.
What was measured	Steps taken, answers marked right, text put in front of the AI, tokens and seconds.	**The controls**	Nothing at all. The documents alone. The folder present but not loaded. A folder with its map made wrong on purpose.
The sample	10 questions, asked 3 times each, in 19 set-ups: 570 sessions. The 19 are every mix of the 4 instruction files loaded or not (16), the documents alone, nothing, and a folder with a wrong map. Then 10 more questions held back as a check.	**The decision rule**	On steps, a difference counts only if it shows on 9 or 10 of the 10 questions. Any fall in right answers, or any rise in text, fails the folder. Written down before any document existed.

How bias was kept out

Read down and follow the arrows. Each step sits under whoever did it. On the right is what that step is there to shut out.

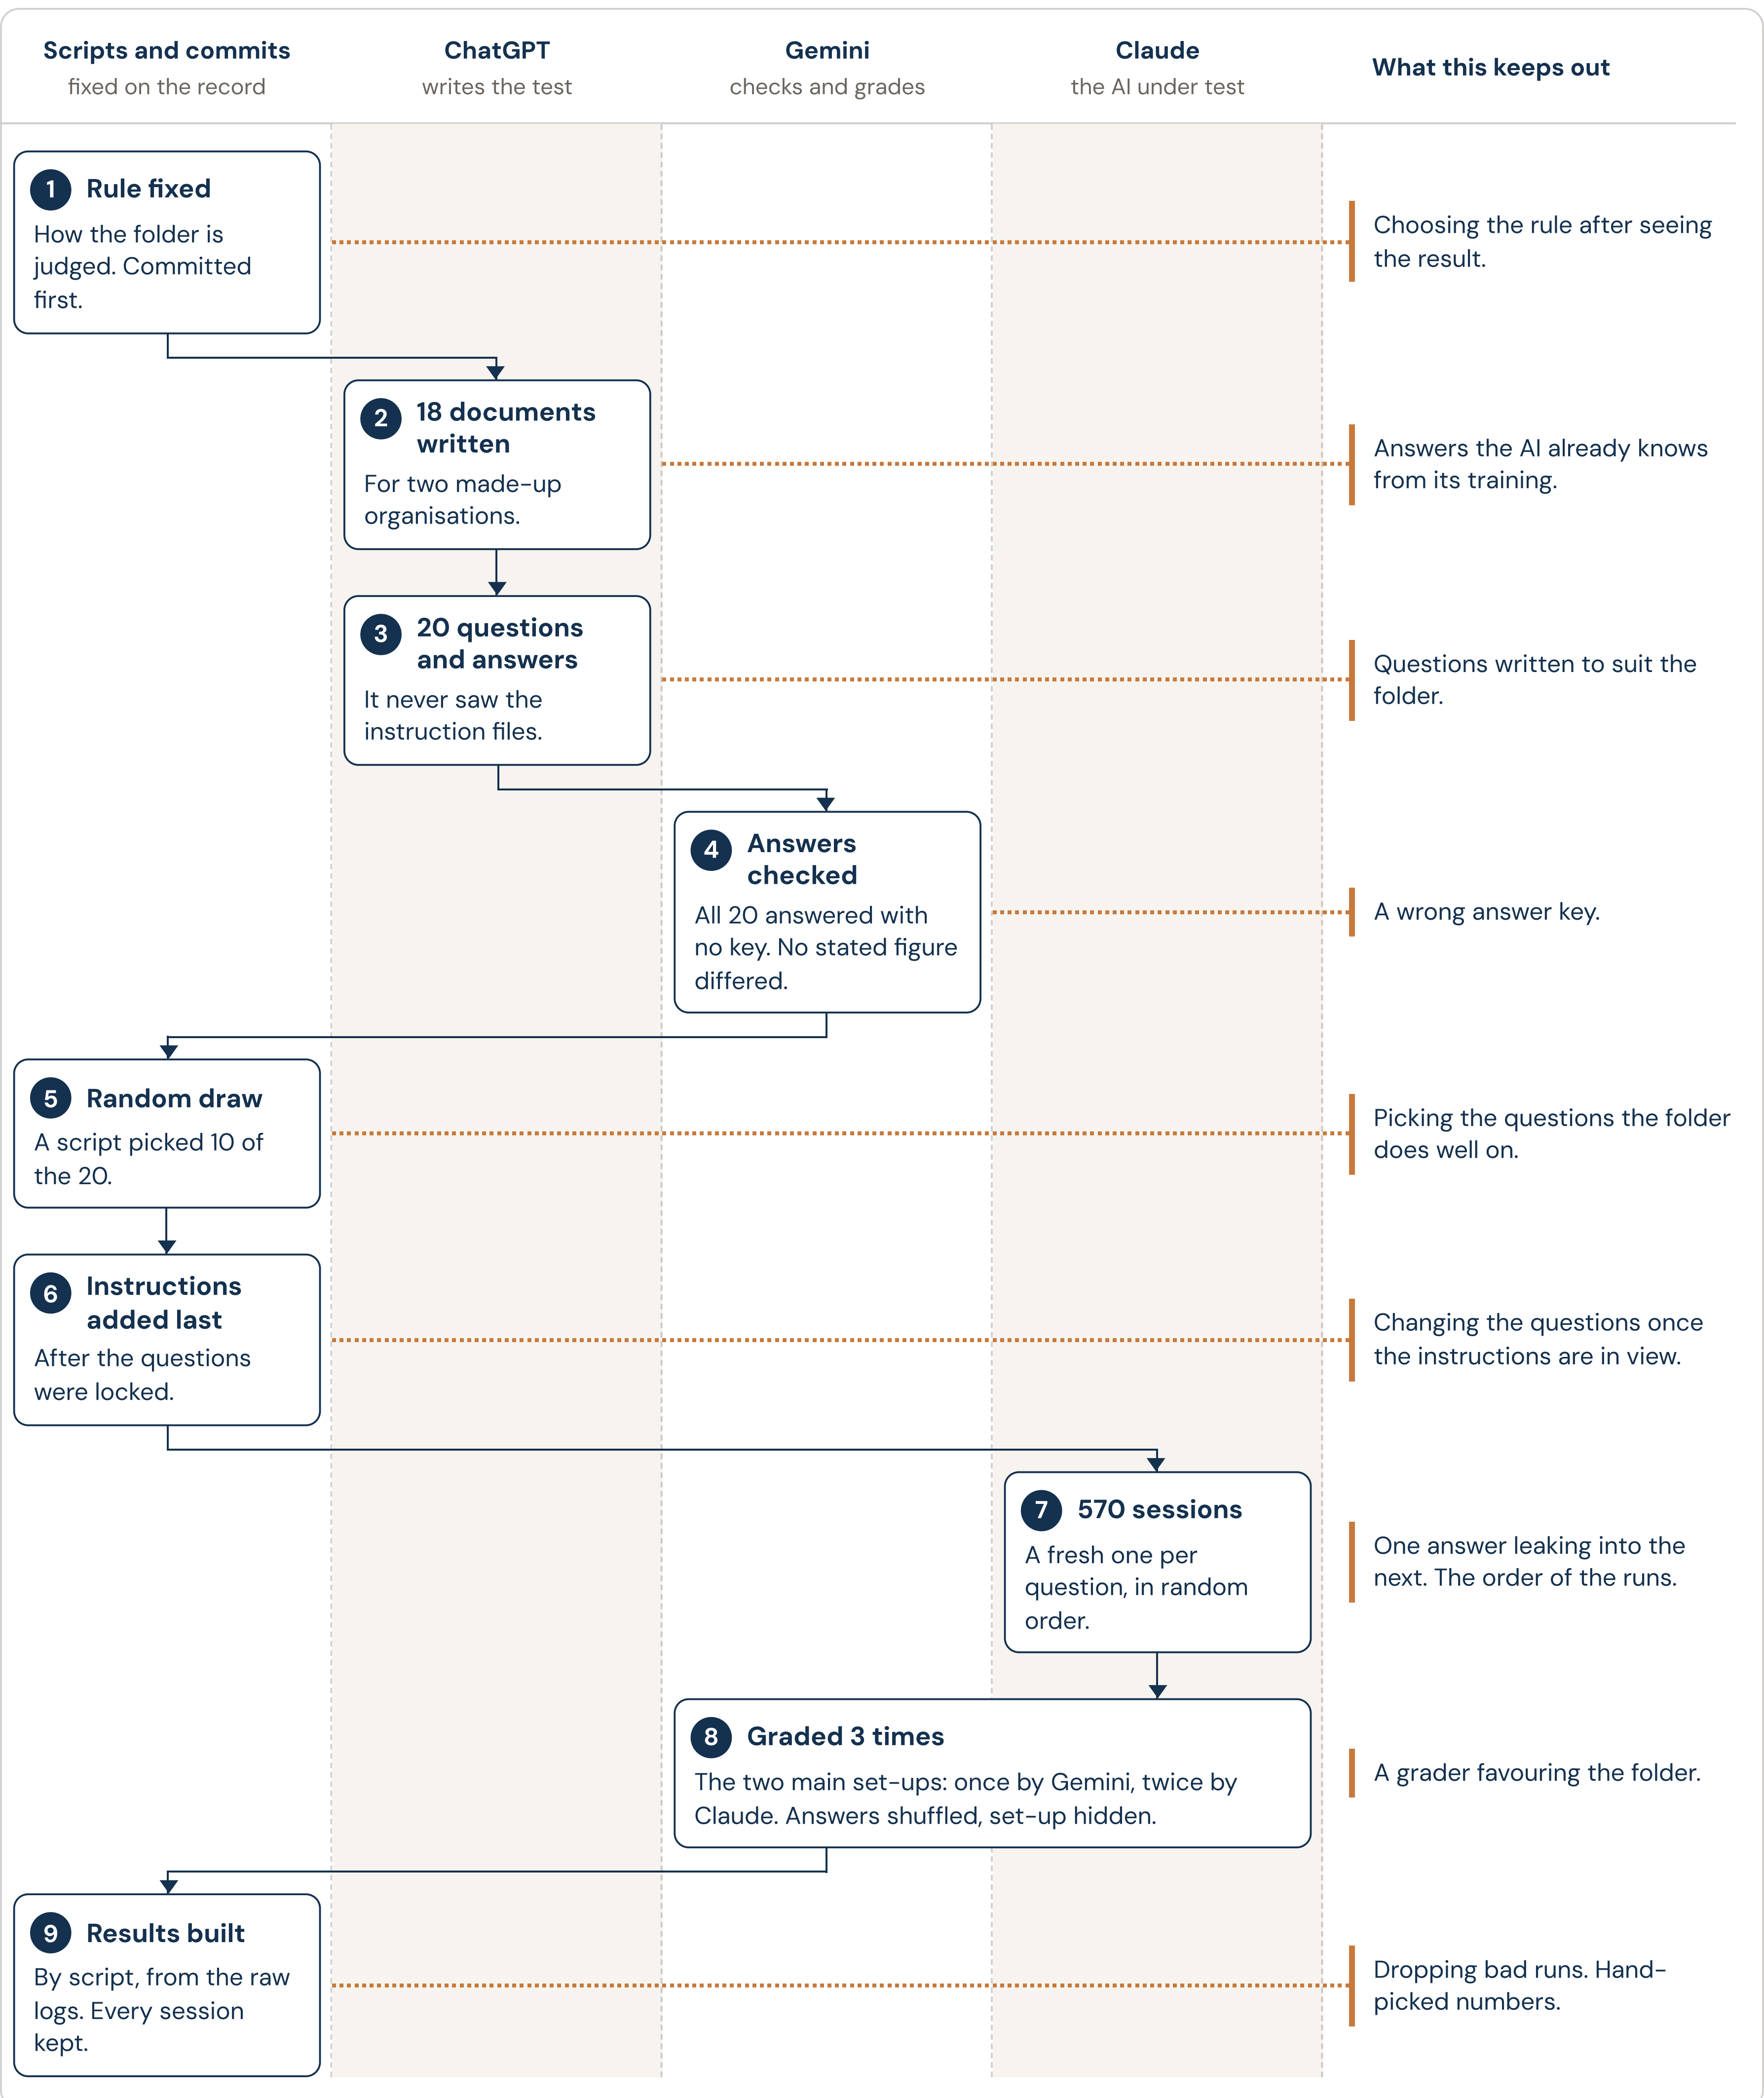
Scripts and commits
fixed on the record
ChatGPT
writes the test
Gemini
checks and grades
Claude
the AI under test
What this keeps out
1 Rule fixed
How the folder is judged. Committed first.
Choosing the rule after seeing the result.
2 18 documents written
For two made-up organisations.
Answers the AI already knows from its training.
3 20 questions and answers
It never saw the instruction files.
Questions written to suit the folder.
4 Answers checked
All 20 answered with no key. No stated figure differed.
A wrong answer key.
5 Random draw
A script picked 10 of the 20.
Picking the questions the folder does well on.
6 Instructions added last
After the questions were locked.
Changing the questions once the instructions are in view.
7 570 sessions
A fresh one per question, in random order.
One answer leaking into the next. The order of the runs.
8 Graded 3 times
The two main set-ups: once by Gemini, twice by Claude. Answers shuffled, set-up hidden.
A grader favouring the folder.
9 Results built
By script, from the raw logs. Every session kept.
Dropping bad runs. Hand-picked numbers.

Where it fell short: the set-up was hidden from the graders, but they could often tell the two apart from how the answers were written. Claude both answered and graded. The builder wrote the prompts for the other AIs and the folder itself.

Where to look. `START-HERE.md` is the guide. `results.md` has every figure, tie and loss. `folder-change/` has the follow-up.

Weekly Comp #14, The Control. Run on 12 October 2026, with follow-ups on the 13th.