AI the vL Way
Less hype. More help.

Does the folder earn its place?

A folder of instructions is meant to help an AI answer questions about a set of documents. This test checked whether it does.

No. The folder got the AI to its answers in half the steps, but its answers were marked right less often than with the documents alone.

The experiment

The same 10 questions, asked of the same AI, in three set-ups. Each set-up was asked every question 3 times.

Nothing
No documents, no folder. The AI only has the question.
No instructions
0 marked right, of 30
2.00 steps per question

Documents alone
18 documents to search. No instructions.
No instructions
30 marked right, of 30
5.03 steps per question

Documents with the folder
The same 18 documents, plus 4 instruction files read first.
map · role · rules · examples
23 to 28 marked right, of 30
2.37 steps per question

The brief asks for the first comparison: the folder against nothing. The folder wins that one, because the documents do the work. The middle set-up is the control that matters. It shows what the 4 instruction files add on top of the documents.

The results

🟧 Documents with the folder ■ Documents alone

Steps to reach an answer

Searches and files opened, per question. Fewer with the folder on 10 of 10.

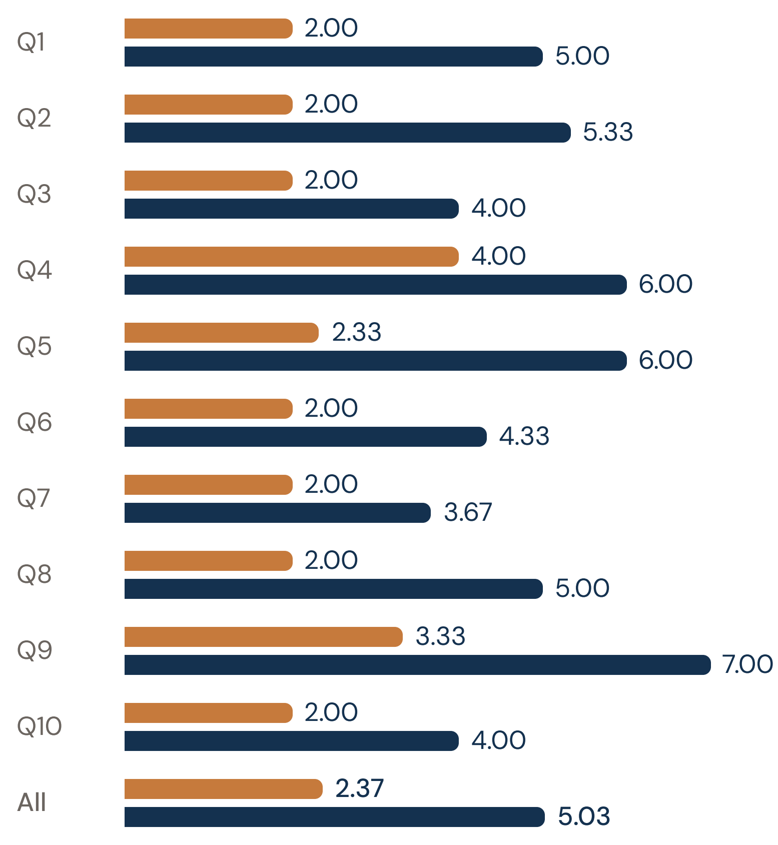

Answers marked right, of 30

Each of the three gradings.

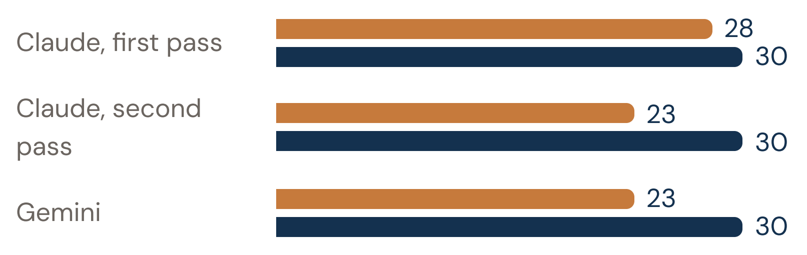

Text put in front of the AI

Characters per question. The lighter part is the folder's own files, read every time.

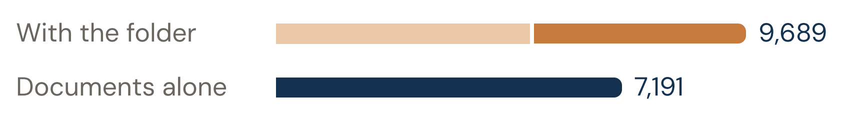

Cost per question

Tokens are the units an AI is billed in. The whole conversation is sent again at every step, so fewer steps means fewer tokens in, even with more text.

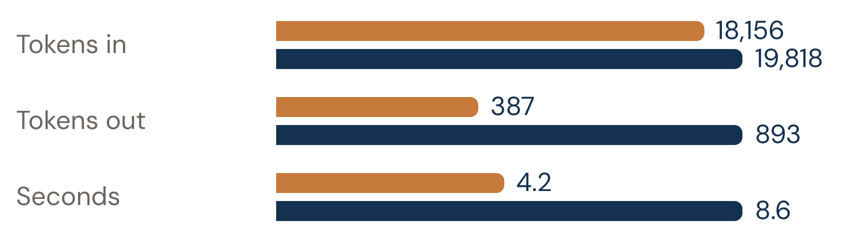

The rule, fixed before the run

The folder earns its place only if all three hold.

✓ **Steps clearly fall**
Met. Fewer on 10 of 10 questions.

✗ **Right answers do not fall**
Not met. 23 to 28 against 30.

✗ **Text does not rise**
Not met. 9,689 against 7,191.

Why the folder lost marks

Our reading, which is not proved. The folder's rules and examples push for short answers. Its answers were right on the main point and less than half the length. The answer key also wanted the detail a longer answer includes, such as how far over a limit something was.

The follow-up: one line, tested on a copy

After the verdict, one line was added to a copy of the folder's rules: "Say what would otherwise apply, and the difference." Then the three set-ups were run again, on the 10 exam questions and on the 10 kept aside.

✓ **The lost marks came back**
30 of 30 on all three gradings. It was 23 to 28 on the exam. The 10 kept aside were already 30 of 30, and stayed there.

✓ **The steps did not move**
2.33 against 5.20. Fewer on 10 of 10.

✗ **The text is still more**
9,579 against 7,665. By the same rule, still not earned.

The line was written after seeing which answers lost marks, so it is fitted to those questions. The first verdict stands.

What held up

- Most of the saving on steps comes from the map: 2.50 with the map alone, 6.07 with nothing loaded.
- A re-run on a second computer, set up from nothing, gave the same result on steps: 2.30 against 5.20.
- On 10 questions kept aside: fewer steps on all 10, and 30 of 30 marked right in both set-ups. The loss of marks did not repeat.

What this does not show

- One AI model, made-up documents, ten questions, one day.
- The graders could tell the two set-ups apart from the writing.
- The second computer used the same account. Nobody outside has run it yet.

The design

How the test was built. What makes it an experiment and not a demonstration.

The one thing changed	Whether the folder's 4 instruction files are loaded before the question.	**Kept the same**	The AI model and its settings, the 18 documents, the 10 questions, the tools (search, list, open), and a fresh session for every question.
What was measured	Steps taken, answers marked right, text put in front of the AI, tokens and seconds.	**The controls**	Nothing at all. The documents alone. The folder present but not loaded. A folder with its map made wrong on purpose.
The sample	10 questions, asked 3 times each, in 19 set-ups: 570 sessions. The 19 are every mix of the 4 instruction files loaded or not (16), the documents alone, nothing, and a folder with a wrong map. Then 10 more questions held back as a check.	**The decision rule**	A difference counts only if it shows on 9 or 10 of the 10 questions. Written down before any document existed.

How bias was kept out

Read down and follow the arrows. Each step sits under whoever did it, so no one hand touches two jobs that should be apart. On the right is what that step shuts out.

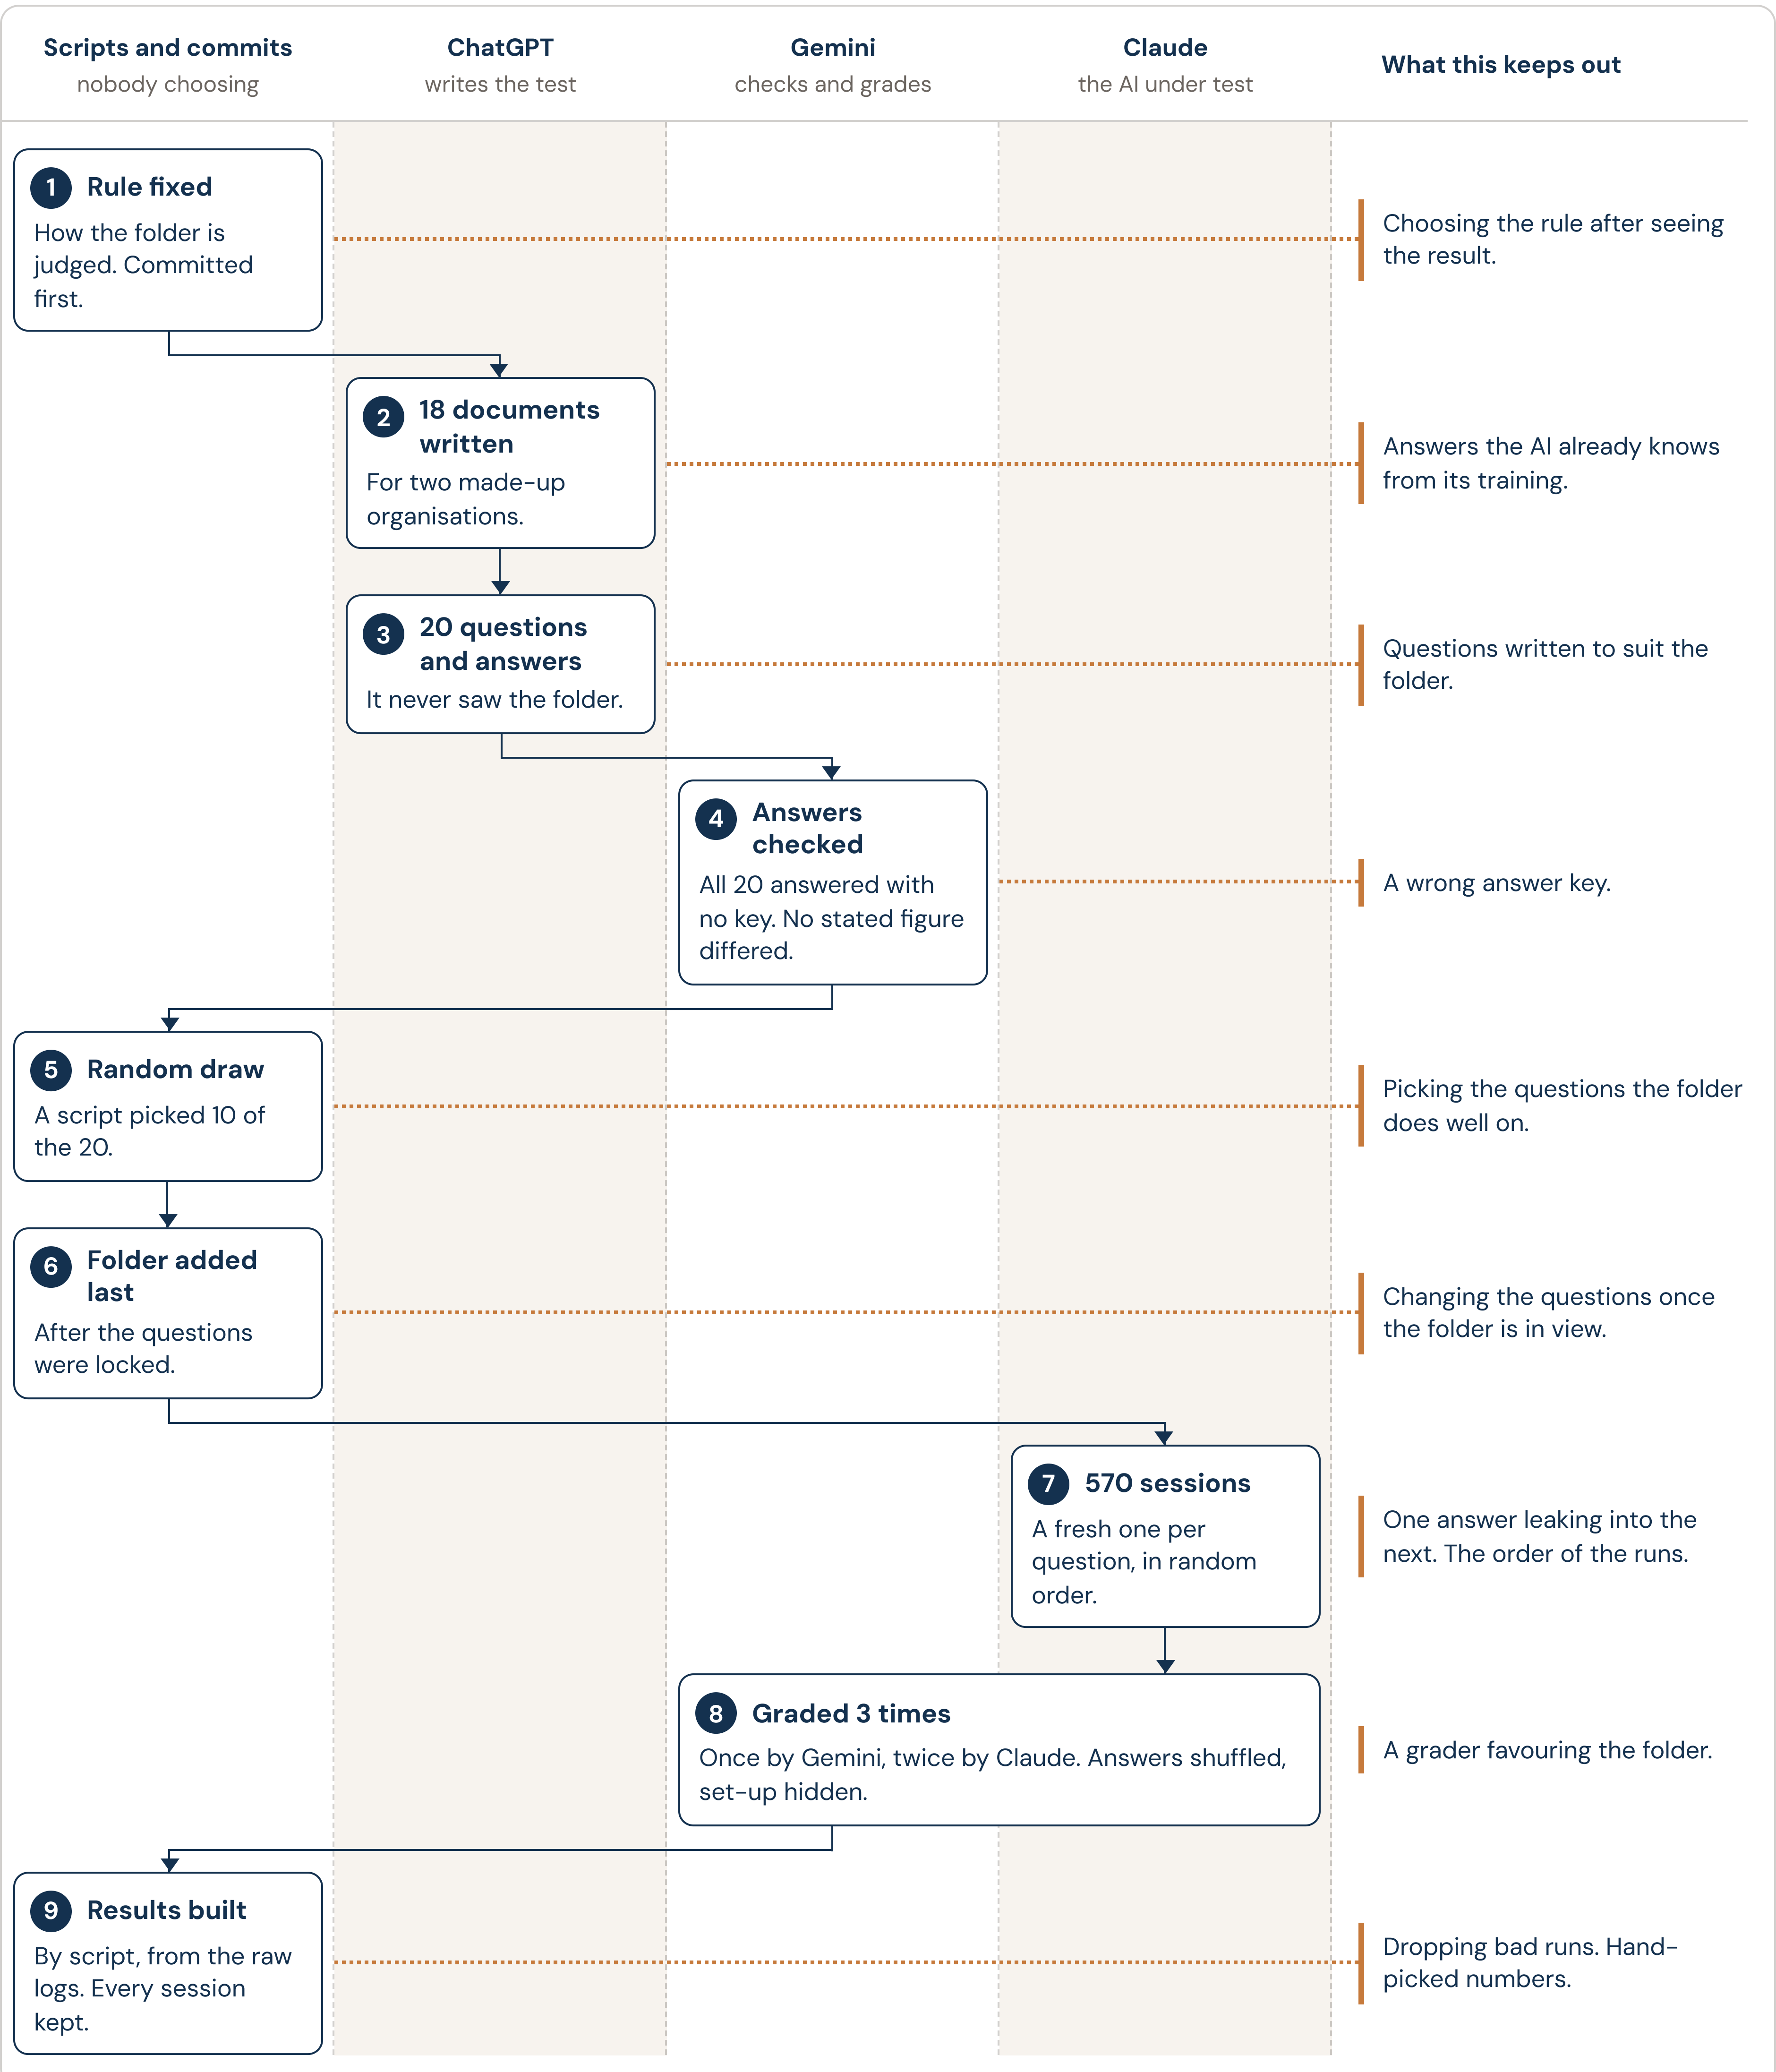

Scripts and commits nobody choosing	ChatGPT writes the test	Gemini checks and grades	Claude the AI under test	What this keeps out
① **Rule fixed** How the folder is judged. Committed first.				Choosing the rule after seeing the result.
	② **18 documents written** For two made-up organisations.			Answers the AI already knows from its training.
	③ **20 questions and answers** It never saw the folder.			Questions written to suit the folder.
		④ **Answers checked** All 20 answered with no key. No stated figure differed.		A wrong answer key.
⑤ **Random draw** A script picked 10 of the 20.				Picking the questions the folder does well on.
⑥ **Folder added last** After the questions were locked.				Changing the questions once the folder is in view.
			⑦ **570 sessions** A fresh one per question, in random order.	One answer leaking into the next. The order of the runs.
		⑧ **Graded 3 times** Once by Gemini, twice by Claude. Answers shuffled, set-up hidden.		A grader favouring the folder.
⑨ **Results built** By script, from the raw logs. Every session kept.				Dropping bad runs. Hand-picked numbers.

Where it fell short: the set-up was hidden from the graders, but they could often tell the two apart from how the answers were written.

Where to look. `START-HERE.md` is the guide. `results.md` has every figure, tie and loss. `folder-change/` has the follow-up.

Weekly Comp #14, The Control. Run on 12 October 2026. AI the vL Way. Less hype. More help.